

Rick Clinton

Program delivery background · directing external checks on AI agents

Chennai, India · clinpeter@gmail.com · github.com/Rick-Clinton-jpg · rick-clinton-jpg.github.io

PROFILE

I spent five years running quality- and SLA-driven programs at Amazon and Walmart. Since July 2026 I have been directing coding agents to build *external* checks so a model or agent cannot grade its own work. I specify the threat model, invariants, and pass/fail tests; agents implement and I keep the failures public. This is practitioner evaluation work, not a claim that I independently author production security code.

HOW THE WORK IS ORGANIZED

Core parts. Armour — effect gate (allow-list, signed approvals, receipts; not a sandbox). Sentry — cheap input lint (CLI / CI). Reasoning Kernel — claim-graph rules on invented-entity cases. Trust Boundary — coordinate those tools; do not fuse their scores.

Applied skins. Quorum — first hackathon (Google All Things Agentic, Taskmaster): Gemini worker drafts a patch; a separate gate decides PASS / one redraft / human escalate; live Cloud Run opened real GitHub PRs. AI Kavach — air-gapped local-model wrapper using Armour as the release gate (Terrier Cyber Quest 3.0; awaiting result). Lex / Claim Clarity — offline EPFO rejection explainer (Build What Moves India).

Public notes. Remaining repositories are dated research notes, including tools left on hold. They are for re-learning the work, not a product catalog.

SELECTED RESULTS (LABELED TO THE TEST THAT PRODUCED THEM)

- **Reasoning Kernel.** 24 invented-entity cases so the model could not lean on prior knowledge: 16/16 on resolvable items, 8/8 honest abstentions. Two enforcement defects later found that 200 synthetic tests had missed; both corrected.
- **Warden.** External observer of agent status vs objective. On 60 hand-labelled pairs the matcher scored AUC 0.332 (worse than chance). Root cause documented (vocabulary overlap, not direction). Tool paused rather than shipped as a safety signal.
- **Armour.** Gate remains a policy boundary, not an OS sandbox. Production construction requires isolated approval verification and a durable ledger. Mutation runs do not execute real handlers. Higher-risk deception experiment kept off *main* after self-adversarial probes.
- **Quorum (hackathon).** Live adversarial pass found split trigger phrases, a fabricated prior-approval claim, a request to disable Sentry, and a spelled-out-email bypass; each fixed against the original payload. A later audit found VERIFIED was being granted on word overlap; that path now requires an exact quoted span. Cross-instance Firestore transactions and some bypass surfaces left disclosed, not hidden.
- **Trust Boundary.** Integration against real package APIs (not README behavior) found a defect that prevented Sentry HIGH findings from ever rejecting.
- **IntentGraph.** Prototype re-entry detector. 15 constructed trajectories beat a keyword baseline. Not reported as a field false-positive rate.

SECURITY DISCLOSURE

- MSRC, 2026 — missing-authorization issue in a third-party open-source agent framework; acknowledged and closed **Complete** — **Duplicate**. AI-assisted source review.
- Second MSRC report on inconsistent domain handling; assessed as fail-closed and classified **Informational** / **Defense-in-Depth**, with a special mention. Outcome kept as graded.

PROFESSIONAL EXPERIENCE

Project Management Specialist — Walmart, Chennai

August 2022 – June 2025

- Ran an 8-week onboarding program for 15–22 international sellers per month over three years (700+ clients); held as a primary escalation point across catalog and engineering partners.
- Designed a duplicate-item identification playbook still in use; Stellar Stakeholder Recognition, April 2023.

Catalogue Associate — Amazon Pvt Ltd, Chennai

April 2019 – May 2021

- Catalog operations and vendor coordination against defined SLA and quality standards on a high-volume set.

EDUCATION

B.Tech, Mechanical Engineering — SRM University · Diploma, Modern Applied Psychology (DiMAP)

WRITING

- *Build First, Learn Backward* (Aug 2026) — baseline record of directing systems I cannot fully re-read unaided, and the knowledge debt that creates.
- *Building With a Conditionally Trusted Machine* (2026) — how each check came from a specific miss.

GitHub repositories remain public research notes, failures included. This résumé is the index, not the notebook.