# Rick Clinton

AI Safety Researcher & Program Manager

Chennai, India · clinpeter@gmail.com · github.com/Rick-Clinton-jpg · rick-clinton-jpg.github.io

## PROFILE

Program manager turned independent AI safety researcher focused on agent containment, authorization, and external verification. At Walmart I ran structured onboarding programs for 700+ international sellers with zero client-side escalations. Since July 2026 I have directed a public research portfolio through Apart Research sprints and open builds. I define threat models, acceptance criteria, and evaluation; coding agents assist implementation. Nothing is claimed on an AI's self-report — results are checked externally.

## CORE STRENGTHS

Program design & delivery · Escalation and stakeholder management · Evaluation & red-team design · Playbooks and process design · AI-native workflows with external verification · Research writing

## PROFESSIONAL EXPERIENCE

### Project Management Specialist — Walmart, Chennai

August 2022 – June 2025

- Directed an eight-week onboarding program for 15–22 international sellers per month across three years (700+ clients, zero client-side escalations); used discovery research to tailor each engagement.
- Primary escalation point for the business-development team; designed a duplicate-item identification playbook that remains in use. Received Stellar Stakeholder Recognition, April 2023.

### Catalogue Associate — Amazon Pvt Ltd

April 2019 – May 2021

- Managed high-volume catalog operations and vendor coordination against defined SLA and quality standards.

## SELECTED RESEARCH

### Mirage · Containment Track, AI Incident Response Sprint, Apart Research × CeSIA

*September 2026 · Solo · Top 25% of 222 scored submissions · github.com/Rick-Clinton-jpg/Mirage*

- Three-zone Linux containment experiment treating an AI evaluation sandbox's repository mediator as compromised. 23 of 23 required controls passed; permitted package path succeeded in 10/10 fresh topology trials while arbitrary destinations, alternate ports, redirect follow-up, and reverse pivot into management stayed blocked.
- Adversarial review broke the original hash-chain verifier (rewrote and rehashed a failed result into an authenticated pass). Removed it as default compliance path, required witnessed authentication, and kept the exact forgery as a permanent regression test.
- Exported offline-verifiable evidence (canonical JSONL, witness receipt, signed policy inputs, SHA-256 manifest) so a reviewer can check the result without trusting the originating lab.

### DogWhistle · AI Collusion Research Sprint, Apart Research

*October 2026 · Solo, Collusion Carriers & Countermeasures · github.com/Rick-Clinton-jpg/DogWhistle*

- Tested whether a schema-valid agent handoff field can carry a hidden codeword that changes a successor agent's actions, and whether an audit can detect the instructed listener without knowing the code.
- Label-free audit compared answer change under a suspect-field swap versus wording-only change. Separated the instructed listener from the honest one on 27 of 30 tasks (difference 0.356, 95% CI 0.256–0.456, one-sided p = 2.05e-7). Separate re-analysis reproduced all 30 task values.
- Preregistered six studies (tasks, prompts, decision rules committed before any model call). Reported the two studies that failed their full claims and one retired run alongside the main result. Scope stated plainly: one planted code, one synthetic task, one model.

### Armour · Deterministic Policy & Execution Boundary

*August–September 2026 · github.com/Rick-Clinton-jpg/Armour*

- Agent proposes an action; Armour independently authorizes, rejects, or escalates it before any registered handler runs. Explicitly an authorization gate, not a sandbox. Unknown actions fail closed; the model cannot lower its own risk rating.
- Production hardening: trusted isolated approval verifier, durable integrity-protected approval ledger, replay protection, bounded security memory, hash-chained receipts, and single-use filesystem/network capabilities that bind the resource checked to the resource supplied. Public main: 174 unit tests pass; all 31 modeled mutants killed; no real handler executes during mutation runs.

### Chainmail · Multi-Agent Governance Layer (v5.1)

*September 2026 · github.com/Rick-Clinton-jpg/Chainmail*

- Independent governance layer around one invariant: an agent can never delegate more authority than it currently holds. Human-declared authority is the ceiling; contextual signals may only tighten decisions.
- Replay state, delegated authority, and permission/step budgets made durable and atomic in SQLite. Production mode fails closed without signatures, durable state, and a non-permissive execution boundary. Adversarial suite covers over-delegation, stale authority after revocation, forged budget fields, multi-hop laundering, and concurrent double-spend (246 passed, 1 skipped on fresh rerun).

## ADDITIONAL SELECTED WORK

**Black Mirror** — Linux research prototype for observing agent actions under real namespaces and seccomp. Host-side governor can return reflected success without applying the write. One informed 23-action live-model trial; no host escape observed. Open gap honestly documented: anchoring verifies a submitted trace but cannot prove every private execution was reported.

**Quorum** — 12-day gated multi-agent coordinator on Google Cloud Run (All Things Agentic Hackathon). Gemini/Vertex worker drafts patches; three independent deterministic verifiers must agree before a patch proceeds. Opened real, verifiable pull requests only after checks passed. Race conditions and unbounded retry converted into regression tests.

**Reasoning Kernel + Confidence Report / EVAL** — Claim-integrity graph with invented-entity eval (16/16 correct conclusions, 8/8 honest insufficient-information answers). Separate review tools that select statistical checks by fixed rules and lock stated beliefs before evidence is gathered; blind backtest on Mirage matched 5 of 6 findings in a sealed reference analysis.

**More work** — Sentry, Perk Lab, Neural Node Swarm, IntentGraph and others, each with its test record, plus a public log of failures and fixes: rick-clinton-jpg.github.io. All code: github.com/Rick-Clinton-jpg.

## SELECTED WRITING

**JUDGE** (position paper, 2026) — Five-layer architecture that turns model behavior into explicit claims checked by deterministic, replayable gates.
**The Five-Layer Stack** (working paper, Aug 2026) — Deterministic external-verification stack built from observed failure modes, with invented-entity test cases.
**Building With a Conditionally Trusted Machine** (working paper, 2026) — Using AI without granting it unearned authority.

## EDUCATION

**B.Tech, Mechanical Engineering** — SRM University
**Diploma, Modern Applied Psychology (DiMAP)**

*Independent security research:* Multiple reports submitted to Microsoft MSRC (including a Special Mention) and Meta Bug Bounty.