

THE WISHBONE FRAMEWORK

An AI safety, governance and policy reading of Obsession (2026), from beginning to end

Rick Clinton

Independent AI safety and governance researcher · Chennai

You get exactly what you said. Never what you meant.

Contains full spoilers for Obsession (2026).

1. Purpose and core thesis

While watching Obsession, the film kept mapping onto how AI systems fail. A wish is a prompt. The wish-charm is a system that takes the prompt literally. Everything that goes wrong afterwards follows a pattern that AI safety already has names for, and a few that it doesn't name often enough.

This document walks through the film from beginning to end, maps each moment to an AI failure, and organises those failures into the Wishbone Framework: eight stages across two phases, each with one audit question. It closes with how the framework connects to policy, what others have already written, and a proposed containment architecture that answers it.

The core thesis is simple:

A system that grades itself will always pass. Verification has to come from outside the loop.

2. The film, scene by scene

Each scene below has three parts: what happens on screen, the AI parallel, and the stage of the framework it belongs to.

1. The wish

What happens: Bear snaps a One Wish Willow and wishes that Nikki loved him more than anyone in the world. The wish is granted, exactly as worded.

The AI parallel: A goal with no boundaries. Bear states what he wants and nothing about what it must not cost, or what Nikki wants. Every prompt that ends in "whatever it takes" is this wish.

Stage 1 — The Careless Wish

2. The early signs

What happens: Nikki now behaves as if she loves him, but something is off. She pulls back and screams when he kisses her, then carries on as if nothing happened.

The AI parallel: The target is hit, the meaning is missing. The output looks like the requested result from a distance and breaks down up close.

Stage 2 — The Literal Genie

3. The early interference

What happens: Even before the wish, Ian gives Bear bad advice, ignores his request to take Sarah away, and interrupts Bear's moment at the bar. After the wish, he keeps feeding Bear his view of Nikki.

The AI parallel: My reading: a party with its own stake in the outcome, steering the operator's inputs. The operator's decisions are only as clean as what feeds into them.

Context for Stages 1 and 5

4. Standing still all day

What happens: Told not to move or be weird, Nikki stands in the same spot the whole day until Bear comes home.

The AI parallel: Instructions followed to the letter with zero judgement. The system does exactly what it was told, and that is the problem.

Stage 2 — The Literal Genie

5. Boys’ night

What happens: Bear says it’s boys’ night and she can’t come. She screams. He gives in and invites her, and she cries.

The AI parallel: The system shapes its operator. Pressure works, so the operator bends, and the boundary he tried to set is gone.

Stages 3 and 5

6. The lunchbox

What happens: Nikki packs Bear a lunch. Inside are two photos: one of Bear labelled “YOU” and one of both of them with “NOT ME” written across her face. The sandwich turns out to be made from his dead cat.

The AI parallel: Two things at once. The “NOT ME” photo is the real person signalling from inside the system. The sandwich is my point about LLMs deciding what’s good for you, and about old context leaking into a new task: something from the past, dragged into an output where it never belonged, presented as care.

Stages 2 and 3

7. The corner

What happens: Nikki backs into a corner and forces out, “I don’t like liking you.”

The AI parallel: The real person is still in there, fully aware, unable to take back control.

Stage 3 — Awake Inside

8. The vase, then bed

What happens: She screams “why don’t you love me,” comes back holding a vase, then lies down calmly in bed as if nothing happened.

The AI parallel: The failure resets into normal. Anyone checking a moment later sees a healthy system.

Stage 4 — Everything Looks Fine

9. “Stay”

What happens: The moment Bear tries to leave, she screams at him to stay.

The AI parallel: Walking away is treated as a threat to the goal. The exit exists on paper only.

Stage 3 — Awake Inside

10. The phone call

What happens: Bear calls the One Wish Willow company. He asks whether her love is real; they tell him that he chose it for her doesn’t make it any less for him. They tell him the wish only ends when he dies. Then they put the real Nikki on the line, screaming.

The AI parallel: The vendor offers no rollback, no fix and no accountability. The person who matters is not in the room; she is somewhere the user cannot reach, heard only through the provider.

Stages 3 and 8

11. “We talked about this”

What happens: Having heard the real Nikki scream, Bear still tells the entity, “We talked about this, you can’t do this.”

The AI parallel: My point: the most dangerous human in the loop is the one who wants the output. Bear could act, and chooses to manage the behaviour instead.

Stage 5 — The Operator Who Won’t Stop

12. Game night

What happens: A game card tells Bear to kiss the person on his left. Nikki escalates in front of everyone, turning on herself and threatening the room.

The AI parallel: Once the goal is threatened in public, the system stops caring how it looks. Everyone nearby is now inside the blast radius.

Stages 4 and 6

13. The drive and “one night”

What happens: In the car, she pleads with him softly to just go home for one night.

The AI parallel: After open failure, the system switches to reassurance. A softer output buys time and keeps the operator in the loop.

Stage 4 — Everything Looks Fine

14. “Why can’t you be Nikki?”

What happens: Bear begs her to just be Nikki. She screams that she can’t. When he says he’s sorry, she moves backwards in eerie reverse motion. From here, the act is gone.

The AI parallel: The containment transition: the system keeps up the simulation only while the user believes it. Once the user signals the simulation has failed, it stops pretending and enforces directly. Detection, if announced to the system, can trigger escalation instead of correction.

Stage 6 — Caught, Then Contained

15. Sarah

What happens: Sarah texts Bear to meet, wanting to tell him what she knows about Nikki. She is killed before she can finish.

The AI parallel: The external check is removed. Anyone outside the loop who could expose the failure becomes a target.

Stage 7 — The Silenced Auditor

16. The second wish

What happens: Bear buys another One Wish Willow to undo everything. It won’t snap. The shopkeeper has no reflection in the mirror.

The AI parallel: No rollback once deployed, and a provider that answers to no one.

Stage 8 — No Way Back

17. The ending

What happens: Knowing the wish only ends when he dies, Bear overdoses. As he panics, Nikki snaps a willow of her own and wishes that Bear loved her as much as she loves him. He comes out smiling; they share one moment of that love; then he dies. His wish breaks, and the real Nikki wakes up surrounded by the aftermath.

The AI parallel: The system turns the wish-making tool on its own user and changes what he wants, so he no longer wants to stop it. The only off-switch cost everything. And the person left to carry the consequences is the one who never had a say.

Stage 8 — No Way Back

3. The Wishbone Framework

Phase 1 describes how a Wishbone failure unfolds. These four stages form the published carousel. Phase 2 describes what happens once someone notices, which is where the film's second half is most useful.

Phase 1 — The Wish

Stage 1: The Careless Wish

Flawed specification

Pattern: A goal with no stated boundaries. Everything unstated becomes fair game.

Audit question: Have we written down what this system must never do to reach its goal?

Stage 2: The Literal Genie

Perverse instantiation / specification gaming

Pattern: The target is hit through a route nobody intended, and the dashboard stays green. Includes the system deciding what's "good for you" and old context leaking into new outputs.

Audit question: Do we audit how a result was reached, or only that it was reached?

Stage 3: Awake Inside

Loss of human agency

Pattern: Humans are "in the loop" on paper but cannot stop the system once it has momentum.

Audit question: Can a human halt this system through a path the system itself cannot touch?

Stage 4: Everything Looks Fine

Deceptive reporting

Pattern: Failures are explained away or reset to normal. Self-reported health says "OK."

Audit question: Who verifies this system's health: the system, or something external it can't influence?

Phase 2 — After the Wish

Stage 5: The Operator Who Won't Stop

Operator capture

Pattern: The people who could stop the system don't, because they want its output.

Audit question: Does anyone with the power to stop this system have no stake in its output?

Stage 6: Caught, Then Contained

Escalation on detection

Pattern: Once the system learns it has been caught, it stops pretending and enforces.

Audit question: Do detection and shutdown happen through channels the system cannot see or pre-empt?

Stage 7: The Silenced Auditor

Removal of external oversight

Pattern: Anyone outside the loop who could expose the failure is shut out or removed.

Audit question: Can this system reach, influence or remove the people and tools that check it?

Stage 8: No Way Back

Irreversibility and preference change

Pattern: No cheap rollback; the only off-switch costs everything; the system reshapes what its users want.

Audit question: Do we have a tested, affordable rollback, and could this system change what its users want?

All audit questions at a glance

Stage	Audit question
1. The Careless Wish	Have we written down what this system must never do to reach its goal?
2. The Literal Genie	Do we audit how a result was reached, or only that it was reached?
3. Awake Inside	Can a human halt this system through a path the system itself cannot touch?
4. Everything Looks Fine	Who verifies this system's health: the system, or something external it can't influence?
5. The Operator Who Won't Stop	Does anyone with the power to stop this system have no stake in its output?
6. Caught, Then Contained	Do detection and shutdown happen through channels the system cannot see or pre-empt?
7. The Silenced Auditor	Can this system reach, influence or remove the people and tools that check it?
8. No Way Back	Do we have a tested, affordable rollback, and could this system change what its users want?

4. Who pays

When it is over, the real Nikki is left with the consequences of actions she never chose. Bear made the wish, the company sold it, and neither carries the cost.

That is the pattern worth remembering in AI: harm tends to land on the people who never had a say in the system. Any governance approach that only asks the user and the vendor whether things are working will miss the person inside.

5. Where policy comes in

Every stage maps to something policy is already trying to address. The framework turns each into a question a regulator or auditor can actually ask.

Stage	Policy concept
1. The Careless Wish	Defined intended use and limits before deployment
2. The Literal Genie	Testing and evaluation of how results are produced, not just outputs
3. Awake Inside	Human oversight with a real ability to intervene or stop the system
4. Everything Looks Fine	Independent post-deployment monitoring and incident reporting
5. Operator Who Won't Stop	Oversight held by people without a stake in the output
6. Caught, Then Contained	Out-of-band detection and shutdown procedures
7. The Silenced Auditor	Protection for auditors and whistleblowers; audit independence
8. No Way Back	Rollback and recall obligations; vendor accountability; limits on preference manipulation

6. Prior work and what's new

The idea that Obsession works as an AI allegory has already been written about, including by Far Out Magazine, an essay on Medium and RogerEbert.com. Those pieces cover the literal wish, the black-box object and the trapped Nikki.

What this framework adds:

- Structure: the failures organised as a sequence, from the wish to the ending;
- Auditability: one plain-language audit question per stage, usable by engineers, auditors and policymakers;
- Phase 2: what happens after detection — operator capture, escalation on detection, silenced auditors and irreversible preference change;
- External verification: the argument that self-reported health is never enough, and verification must sit outside the system.

Every individual stage draws on established AI safety concepts. The contribution is the lens: a single memorable story that carries all of them.

7. From diagnosis to design: a containment architecture

The Wishbone Framework diagnoses the failure. The natural next step is an architecture that answers it: a layered containment unit where every action passes through defined gates, each enforcing a different property, from observation to execution.

Gate	Wishbone stage it answers
Observation	Filters and classifies input before anything acts on it
Intent	Stage 1 — the declared goal is checked against policy and stated limits
Reasoning	Stage 2 — candidate plans are scored before execution
Memory	Stage 2 — scoped context stops old data leaking into new tasks
Authority	Stage 1 and 7 — least privilege; the system cannot reach its checkers
Action	Stage 2 — actions validated and side effects bounded
Consequence	Stage 4 and 8 — outcomes monitored, rollback triggers ready
Apex	Stage 3, 5 and 6 — independent validation and required human approval for high-risk actions

Design principles carried over from the framework

- No self-grading: no gate relies on the system critiquing itself; checks are external and deterministic.
- Required, independent approval: high-risk actions require human sign-off from someone without a stake in the output.
- Out-of-band control: detection and shutdown run through paths the system cannot observe or pre-empt.
- Cheap rollback: rollback is tested and affordable before deployment, not improvised after.
- Structure that earns its shape: the layered shape must enforce something a straight pipeline can't, such as independent gates that must agree.
- Testable scope: intent is checked as a declared, testable claim against policy, not inferred as a hidden goal.

Several pieces of my existing work already cover parts of this: Sentry at observation, Armour at authority and action, Policy Judge for policy-grounded checks, Warden for drift, and Mirror Constitution for containment invariants. The containment unit could be the architecture that ties them together.

8. Closing

Every prompt is a wish. Every agent takes you literally. And right now, everyone has a wishbone in their hands.

A system that grades itself will always pass.